

The future web: HTTP 3 in Rust

Dirkjan Ochtman
Rust Nijmegen, March 2019

<https://github.com/djc/quinn>

About me



@djc



@djco

- <https://dirkjan.ochtman.nl/>
- Rust projects: Askama, tokio-imap
- Software engineer on ING's Distributed Ledger Technology ("blockchain") team

What is HTTP 3?

- Reformulating HTTP (2) on top of the QUIC transport
- Re-envision TCP after ~40 years of use
- QUIC: started at Google in 2012, IETF working group established in 2016, planning for final drafts in Nov 2018
- (SPDY announced 2009; first draft of HTTP 2 in 2012; RFC 7540 published May 2015; SPDY removed in 2016.)
- Expect QUIC v1 and HTTP 3 RFC in Q3 2019

Table of contents

- **QUIC: problem statement, concepts and process**
- Quinn: design, implementation status, future

Head-of-Line blocking

Application

Transport

Network

QUIC goals

- Minimize latency
- Provide multiplexing
- Changes limited to path endpoints
- Support connection migration
- Provide always-secure transport
- Expect rapid distributed development and testing

QUIC documents

- Invariants: things that cannot change in future versions
- Transport: core protocol logic
- Recovery: loss detection and congestion control
- TLS: how to secure the transport payloads
- QPACK: header compression (static & dynamic)
- HTTP: mapping of HTTP (2) semantics to QUIC



© Wendy Carlyle (CC NC-BY-SA 2.0)

Recovery

- Differences with TCP
 - Monotonically increasing packet numbers
 - No renegeing (reduces memory pressure on sender)
 - More ACK ranges (less reliant on timeouts)
 - Explicit correction for delayed acknowledgements

QUIC TLS

- TLS 1.3 handshake messages with a custom extension
 - Extension for announcing transport parameters
- After handshake, encrypt payloads with negotiated secret
 - Application data is not wrapped in TLS records

HTTP 3

- ALPN protocol "h3" negotiated during TLS handshake
- Builds on HTTP 2, differs in order to leverage QUIC
 - Different header compression scheme
 - Different framing mechanism

Interoperability matrix

	h2o/quickly	quant	ngtcp2	mvfst	picoQUIC	winquic	f5	ATS	quiche	lsquic	nginx_quic	AppleQUIC	quic-go	Quinn	P
client ↓	h2o/quickly	quant	ngtcp2	mvfst	picoQUIC	winquic	f5	ATS	quiche	lsquic	nginx_quic	AppleQUIC	quic-go	Quinn	P
h2o/quickly															
quant		VHDCRZS MBUPE	VHDCRZS MBU		VHDCRZS MBU3P	VHDCRZS P	VHDCRS	VHDCRZS M3	VHDCRS	VHDCRZS M3E	VHDCR			VHDCRZS MBUPE	
ngtcp2		VHDCRZS U	VHDCRZS MBU		VHDCRZS MBU	VHDCRZS	VHDCRS B	VHDCRZS MB	VHDCRS	VHDCS				VHDCRZS MBU	
mvfst															
picoQUIC		VHDCRZS MBUP	VHDCRZS MBU		VHDCRZS MBU3P	VHDCRZS P	VHDCRS B	VHDCRZS MB3	VHDCRS	VHDCRZS MB3	VHDCR	VHDC		VHDCRS MBUP	
winquic		VHDCRZS	VHDCRZS		VHDCRZS	VHDCRZS P	VHDCS	VHDCRZS	VHDCRS	VHS				VHDCS	
f5		HDCS	HDCS		HDCS	HDCS	HDCS	HDCS	HDCS	HCS				HDCS	
ATS		VHDCRS M	VHDCRS M		VHDCRS M	VHDCRS	VHDCRS	VHDCRS M	VHDCRS	VHS		VHDC		VHDCRS M	
quiche															
lsquic					VHDCRS 3			VHDCRS 3		VHDCRS 3					
nginx_quic	-	-	-	-	-	-	-	-	-	-	-	-	-	-	
AppleQUIC		VHDCS	VHDCS		VHDCS 3	VHDCS	VHDCS	VHDCS 3	VHDCS	VHDCS 3		VHDCS		VHDCS	
quic-go															
Quinn		VHDCRZS U	VHDCRZS BU		VHDCRZS BU	VHDCRZS	VHDCS	VHDCRZS 3		VHS			VHDCRZS BUP		
Pandora		VHDC													
f5_test		VHDCS B	VHDCS B		VHDCS 3	V	VHDCS B	VHDCS B3	VHDCS	VHDCS B3				V	
Spindump															

Version Negotiation – Handshake – Stream Data – Connection Close – Resumption – 0-RTT – Stateless Retry

Table of contents

- QUIC: problem statement, concepts and process
- **Quinn: design, implementation status, future**

History

- Started Quinn last year in March
 - Subject of a RustFest May 2018 presentation
- Found that Benjamin Saunders (@Ralith) had also started
 - Merged implementations in Q4, rapidly iterated
- By end of 2018, nicely conforming updated impl
- Released 0.2 release on Jan 21

Goals

- Implement the specifications faithfully
- Have clear abstractions & concepts
- Futures-based architecture – no polling
- Leverage the type system
- Good test coverage at different layers (currently ~77%)

quinn: futures-based application API

quinn-h3

**QPACK,
HTTP 3**

quinn-proto: deterministic protocol logic

**Endpoint: dispatching
datagrams to connections**

**Connection: core
per-connection
protocol logic**

**Crypto: TLS and
encryption**

**Packet: packet encoding
and decoding**

**Stream: streams state
machine tracking API**

rustls and *ring*

- Highly secure, relatively easy to work with/hack on
- Added QUIC TLS needs to rustls:
 - Transport parameters extension
 - Handshake message-based interface
- *ring* algorithms are used directly after handshake
- Behind TLS interface traits to make it pluggable

CryptoSession trait

```
pub trait CryptoSession {
    fn alpn_protocol(&self) -> Option<&[u8]>;
    fn early_crypto(&self) -> Option<Crypto>;
    fn early_data_accepted(&self) -> Option<bool>;
    fn is_handshaking(&self) -> bool;
    fn read_handshake(&mut self, buf: &[u8]) -> Result<(), TransportError>;
    fn sni_hostname(&self) -> Option<&str>;
    fn transport_parameters(&self) -> Result<Option<TransportParameters>, TransportError>;
    fn write_handshake(&mut self, buf: &mut Vec<u8>) -> Option<Crypto>;
}
```

futures, tokio, h2

- Took API inspiration from h2 implementation
 - How to create a good futures-based API?
 - How does the API layer interact with library internals?
- Clean layering between API and protocol logic
- Possible future: `tokio::net::quic`

Ecosystem

- 11 others have contributed to Quinn so far
- In the process of improving Quinn, also contributed to:
 - libc (related to ECN support)
 - tokio
 - ***ring***
 - rustls

Implementation status

- Very conformant draft 18 implementation
- Handles 70 MB/s with 128 kb streams (unoptimized)
- Complete QPACK implementation, not much H3 yet
- async/await implementation for testing on nightly branch
- Pace in transport/tls drafts seems to have slowed down

Open projects

- In progress:
 - H3
 - Task per connection
 - Finish pluggable TLS
- Open:
 - Stream prioritization
 - MTU discovery

quiche

- <https://github.com/cloudflare/quiche>
- Started late last year to replace their C++ version
- Uses BoringSSL bindings for TLS
- Low-level protocol API only

Request for contributions

- <https://github.com/djc/quinn>
- Your contribution is much appreciated
- There are a few issues in the repository
- Happy to mentor anyone to help get started